

Term Jobs - Redefined MVP

- Term Jobs - Redefined MVP
0. Why this redefinition exists

1. MVP scope, restated

2. Flow, end to end

3. Module detail

4. Agent mechanics (unchanged from v4/Final Plan)

5. Build order

6. What's deferred to Phase 2 (unchanged set, for reference)

7. Open questions specific to this MVP

Term Jobs - Redefined MVP

Vendor-facing hiring flow: JD intake -> resume analysis -> interview setup

0. Why this redefinition exists

The original v4 MVP scoped 4 agents around **operations** (Screening, Timesheet Review, Recon/Invoice Validation, FAQ-only support) - it assumed hiring itself was manual and automated what happens *after* a hire (time, billing).

The actual product priority is different: automate the **vendor-facing hiring funnel** - JD in, resumes screened, interviews scheduled - and defer time tracking and billing entirely. This document replaces the agent/module scope in v4 and the Final Project Plan with a leaner MVP matched to that priority. Nothing about the underlying stack, module architecture pattern, or guardrail mechanics changes - only *which* modules and agents ship first.

1. MVP scope, restated

In scope:


Module	Why it's in
identity/	Tenancy (client / consultancy), auth, roles - nothing else works without it
requisition/	Holds the JD and requisition record; state machine for requisition lifecycle
candidate/	Resume intake, candidate-to-requisition linking, interview state

In scope - agents:

Agent	Trigger	Reads	Produces	Guardrail
Job Requirement Agent	requisition.created	Raw JD text	Normalized structured role (title, must-have skills, nice-to-have, seniority, location, rate band)	Always human-reviewed - hiring manager edits before publish
Candidate Screening Agent	candidate.linked_to_requisition	Resume, structured role, existing submissions for that req	Fit score, duplicate-submission flags	Human-reviewed by default; auto-shortlist/auto-reject available as a per-tenant opt-in, off by default
Interview Scheduling Agent	Candidate shortlisted (human-approved)	Interviewer availability, candidate availability (via comms)	Proposed interview slots, confirmed booking	Always human-reviewed for the interviewer side; candidate-facing scheduling comms can auto-send once a human has approved the shortlist

Explicitly out of MVP: - worker/ , compliance/ , time_expense/ , billing/ , analytics/ - deferred in full - Vendor Matching Agent, Rate Benchmark Agent, Performance Insight Agent, Compliance Agent - Timesheet Audit Agent, Invoice Validation Agent - Multi-country support (India-only, single country pack) - Auto-execute for anything beyond candidate shortlisting - no auto-hire, no auto-onboarding

2. Flow, end to end



Every checkpoint above is a LangGraph graph pause, persisted to Postgres - same mechanism as v4's original design, just applied to this narrower flow.

3. Module detail

3.1 identity/

- Tenant types: `client` , `consultancy` (platform admin exists but isn't user-facing at MVP)
- RBAC roles: Hiring Manager, Recruiter (consultancy-side), Admin
- RLS from day one, single country pack (India)
- **Open question carried forward:** does one consultancy span multiple client relationships as a single tenant, or a scoped record per relationship? Blocks nothing at this narrow MVP scope (no billing yet), but should be answered before Phase 2 re-introduces billing.

3.2 requisition/

- Domain entity: Requisition (raw JD, structured role once processed, status)
- State machine: `Draft` -> `Structuring` -> `PendingApproval` -> `Published` -> `Closed`
- Job Requirement Agent writes into `Structuring` , human approval moves it to `Published`

3.3 candidate/

- Domain entity: Candidate, Submission (links candidate to a requisition)
- State machine: `Submitted` -> `Screened` -> `Shortlisted` -> `InterviewScheduled` -> `InterviewCompleted` -> `Rejected/Selected`
- Screening Agent writes into `Screened` ; human action moves to `Shortlisted`
- Interview Scheduling Agent writes into `InterviewScheduled`

4. Agent mechanics (unchanged from v4/Final Plan)

All three agents run through the same guardrail loop:

```
budget gate -> build context -> tool loop (<= N calls) -> pydantic schema validation
-> guardrail check (confidence / reversibility) -> persist decision record
-> [approval queue] or [auto-dispatch, only if tenant threshold allows]
```

- **Cost discipline:** 3-tier model routing (small/mid/large), Langfuse OSS tracing, CI eval harness with golden fixtures before any prompt/model change ships.
- **Decision records:** every agent output persisted as an auditable record, regardless of whether it was auto-dispatched or queued for human review.
- **No agent touches `identity/` directly** - consistent with the "Core stays deterministic" rule from the original architecture.

5. Build order

1. `identity/` - tenancy, roles, RLS
2. `requisition/` + **Job Requirement Agent**
3. `candidate/` + **Candidate Screening Agent**
4. **Interview Scheduling Agent** (comms integration - email/WhatsApp)

Each step is independently testable: you can validate the Job Requirement Agent against fixture JDs before Candidate module exists, then validate Screening against fixture resumes once `candidate/` lands, then Interview Scheduling last since it depends on both prior agents' outputs.

6. What's deferred to Phase 2 (unchanged set, for reference)

- `worker/` , `compliance/` , `time_expense/` , `billing/` , `analytics/`
- Vendor Matching Agent, Rate Benchmark Agent, Performance Insight Agent, Compliance Agent, Timesheet Audit Agent, Invoice Validation Agent
- Support Copilot (real RAG version - static FAQ only, if needed at all, at MVP)
- Second country pack
- Second KYC provider, e-signature, payments - none of these are relevant until `worker/` + `billing/` are back in scope

7. Open questions specific to this MVP

Question	Blocks	Notes
What counts as “interviewer availability” input - calendar integration or manual entry?	Interview Scheduling Agent’s context-build step	Needs an answer before that agent’s tool loop can be specified
Does Screening Agent’s auto-shortlist opt-in ship at MVP or stay off entirely until trust is established?	candidate/ guardrail config	Recommend: off for MVP, revisit after eval data exists
Single JD format (free text) or structured intake form as well?	Job Requirement Agent’s input contract	Affects prompt design and schema validation strictness

This document supersedes the module/agent scope in v4 and the Final Project Plan for MVP purposes only. Architecture patterns, stack choices, and guardrail mechanics from those documents remain the source of truth and are unchanged here.